

Workload Distribution

Member 1 — Data Storage & Sampling Engineer

Responsibilities

1. Data File Handling

- Understand and manage the binary file storing all vectors.
- Ensure correct use of memory-mapped files.
- Implement any helper functions needed for efficient row access.

2. Sampling Module

- Write code to randomly sample vectors from the dataset.
- Ensure sampling does not exceed RAM limits.
- Prepare a sampled dataset for k-means training.

3. Support for Index Building

- Build helper functions that allow the index builder (Member 3) to read vectors efficiently.

Your Output / Deliverables

- Efficient reading of individual vectors from disk.
- A working sampling function.
- Documentation explaining the data format and I/O strategy (memory-mapped approach).

Member 2 — Clustering & Centroid Builder

Responsibilities

1. Train k-Means (**MiniBatchKMeans**)

- Run mini-batch k-means on the sampled vectors.
- Experiment with different numbers of clusters.
- Choose the final cluster count (n_clusters).

2. Centroid Management

- Write centroids to disk in a compact format.
- Provide a function for loading centroids during retrieval.

3. Index Directory Setup

- Define the structure where index files (centroids + inverted lists) will be stored.

Your Output / Deliverables

- A fully trained set of centroids saved to disk.
- Python code for training and loading centroids.
- Documentation explaining clustering strategy, sample size choice, and cluster count.

Member 3 — Index Construction & Inverted Lists Engineer

Responsibilities

1. Implement `_build_index()`

- Read every vector using functions from Member 1.
- Find the nearest centroid for each vector (using centroids from Member 2).
- Assign vector ID to the correct centroid.

2. Create Inverted Lists

- One file per cluster.
- Only store vector IDs to keep index size small.
- Organize files efficiently for fast loading in retrieval.

3. Index Size Optimization

- Make sure all inverted list files meet index size constraints.
- Compress or optimize formats if needed.

Your Output / Deliverables

- Correct implementation of `_build_index()`.
- All inverted list files generated and saved to disk.
- Documentation describing index creation and file-organization decisions.

Member 4 — Retrieval Engineer (`retrieve()` + Evaluation)

Responsibilities

1. Implement `retrieve()`

- Normalize query vector.
- Load centroids from disk (allowed in RAM).
- Compute distances to all centroids.
- Select top n_{probe} clusters (e.g., top 3–10).
- Load inverted list files for selected clusters.
- Load only the needed vectors from the main file.
- Compute cosine similarity for candidates.
- Return top-K similar vectors.

2. Performance Optimization

- Make sure retrieval meets RAM and time limits.
- Profile slow steps and improve them.

3. Evaluation

- Run evaluation notebook.
- Test accuracy (recall) and timing.
- Share results with the team to adjust parameters.

Your Output / Deliverables

- A fast, correct `retrieve()` function.
- Passing all constraints in the evaluation notebook.
- Documentation explaining the retrieval algorithm and optimization choices.

 **Ultra-Clear Summary Table**

Member Main Role	What They Produce
1 Data Storage & Sampling	Efficient data access + sampling code
2 Clustering & Centroids	Trained centroids + centroid files on disk
3 Index Builder	Inverted lists + full _build_index()
4 Retrieval & Evaluation	retrieve() implementation + performance testing